

```

1 %% Machine Learning
2 % Lab 7: Neural Network Training
3 % — Handwritten Digits —
4 %{
5 For this exercise, you will teach a neural network to
6 recognize handwritten digits (from 0 to 9). Automated handwritten digit
7 recognition is widely used today — from recognizing zip codes (postal codes)
8 on mail envelopes to recognizing amounts written on bank checks.
9 %}
10
11 %% Initialization
12 clear ; close all; clc
13
14 %% Setup the parameters you will use for this exercise
15 input_layer_size = 400; % 20x20 Input Images of Digits
16 hidden_layer_size = 25; % 25 hidden units
17 num_labels = 10; % 10 labels, from 1 to 10
18 % (note that we have mapped "0" to label 10)
19
20 %% ===== Part 1: Loading and Visualizing Data =====
21 % We start the exercise by first loading and visualizing the dataset.
22 % You will be working with a dataset that contains handwritten digits.
23 %
24
25 % Load Training Data
26 fprintf('Loading and Visualizing Data ...\n')
27
28 load('ex4data1.mat');
29 m = size(X, 1);
30
31 % Randomly select 100 data points to display
32 sel = randperm(size(X, 1));
33 sel = sel(1:100);
34
35 displayData(X(sel, :));
36
37 fprintf('Program paused. Press enter to continue.\n');
38 pause;
39
40
41 %% ===== Part 2: Loading Parameters =====
42 % In this part of the exercise, we load some pre-initialized
43 % neural network parameters.
44
45 fprintf('\nLoading Saved Neural Network Parameters ...\n')
46
47 % Load the weights into variables Theta1 and Theta2
48 load('ex4weights.mat');
49 w1 = Theta1;
50 w2 = Theta2;
51
52 % Unroll parameters
53 nn_params = [w1(:) ; w2(:)];
54
55 %% ===== Part 3: Compute Cost (Feedforward) =====
56 % To the neural network, you should first start by implementing the

```

```

57 % feedforward part of the neural network that returns the cost only. You
58 % should complete the code in nnCostFunction.m to return cost. After
59 % implementing the feedforward to compute the cost, you can verify that
60 % your implementation is correct by verifying that you get the same cost
61 % as us for the fixed debugging parameters.
62 %
63 % We suggest implementing the feedforward cost *without* regularization
64 % first so that it will be easier for you to debug. Later, in part 4, you
65 % will get to implement the regularized cost.
66 %
67 fprintf('\nFeedforward Using Neural Network ...\n')
68
69 % Weight regularization parameter (we set this to 0 here).
70 lambda = 0;
71
72 C = nnCostFunction(nn_params, input_layer_size, hidden_layer_size, ...
73                   num_labels, X, y, lambda);
74
75 fprintf(['Cost at parameters (loaded from ex4weights): %f '...
76         '\n(this value should be about 0.287629)\n'], C);
77
78 fprintf('\nProgram paused. Press enter to continue.\n');
79 pause;
80
81 %% ===== Part 4: Implement Regularization =====
82 % Once your cost function implementation is correct, you should now
83 % continue to implement the regularization with the cost.
84 %
85
86 fprintf('\nChecking Cost Function (w/ Regularization) ... \n')
87
88 % Weight regularization parameter (we set this to 1 here).
89 lambda = 1;
90
91 C = nnCostFunction(nn_params, input_layer_size, hidden_layer_size, ...
92                   num_labels, X, y, lambda);
93
94 fprintf(['Cost at parameters (loaded from ex4weights): %f '...
95         '\n(this value should be about 0.383770)\n'], C);
96
97 fprintf('Program paused. Press enter to continue.\n');
98 pause;
99
100
101 %% ===== Part 5: Sigmoid Gradient =====
102 % Before you start implementing the neural network, you will first
103 % implement the gradient for the sigmoid function. You should complete the
104 % code in the sigmoidGradient.m file.
105 %
106
107 fprintf('\nEvaluating sigmoid gradient...\n')
108
109 g = sigmoidGradient([-1 -0.5 0 0.5 1]);
110 fprintf('Sigmoid gradient evaluated at [-1 -0.5 0 0.5 1]:\n ');
111 fprintf('%f ', g);
112 fprintf('\n\n');
113

```

```

114 fprintf('Program paused. Press enter to continue.\n');
115 pause;
116
117
118 %% ===== Part 6: Initializing Pameters =====
119 % In this part of the exercise, you will be starting to implment a two
120 % layer neural network that classifies digits. You will start by
121 % implementing a function to initialize the weights of the neural network
122 % (randInitializeWeights.m)
123
124 fprintf('\nInitializing Neural Network Parameters ...\n')
125
126 init_w1 = randInitializeWeights(input_layer_size, hidden_layer_size);
127 init_w2 = randInitializeWeights(hidden_layer_size, num_labels);
128
129 % Unroll parameters
130 initial_nn_params = [init_w1(:) ; init_w2(:)];
131
132 %% ===== Part 7: Implement Regularization =====
133 % Once your backpropagation implementation is correct, you should now
134 % continue to implement the regularization with the cost and gradient.
135 %
136
137 fprintf('\nChecking Backpropagation (w/ Regularization)... \n')
138
139 lambda = 3;
140
141 % Also output the costFunction debugging values
142 debug_C = nnCostFunction(nn_params, input_layer_size, ...
143                         hidden_layer_size, num_labels, X, y, lambda);
144
145 fprintf(['\n\nCost at (fixed) debugging parameters (w/ lambda = %f): %f ' ...
146         '\n(for lambda = 3, this value should be about 0.576051)\n\n'], lambda,
147         debug_C);
148
149 fprintf('Program paused. Press enter to continue.\n');
150 pause;
151
152 %% ===== Part 9: Training NN =====
153 % You have now implemented all the code necessary to train a neural
154 % network. To train your neural network, we will now use "fmincg", which
155 % is a function which works similarly to "fminunc". Recall that these
156 % advanced optimizers are able to train our cost functions efficiently as
157 % long as we provide them with the gradient computations.
158 %
159 fprintf('\nTraining Neural Network... \n')
160
161 % After you have completed the assignment, change the MaxIter to a larger
162 % value to see how more training helps.
163 options = optimset('MaxIter', 50);
164
165 % You should also try different values of lambda
166 lambda = 1;
167
168 % Create "short hand" for the cost function to be minimized
169 costFunction = @(p) nnCostFunction(p, ...

```

```

170         input_layer_size, ...
171         hidden_layer_size, ...
172         num_labels, X, y, lambda);
173
174 % Now, costFunction is a function that takes in only one argument (the
175 % neural network parameters)
176 [nn_params, cost] = fmincg(costFunction, initial_nn_params, options);
177
178 % Obtain Theta1 and Theta2 back from nn_params
179 w1 = reshape(nn_params(1:hidden_layer_size * (input_layer_size + 1)), ...
180             hidden_layer_size, (input_layer_size + 1));
181
182 w2 = reshape(nn_params((1 + (hidden_layer_size * (input_layer_size + 1))):end), ...
183             num_labels, (hidden_layer_size + 1));
184
185 fprintf('Program paused. Press enter to continue.\n');
186 pause;
187
188
189 %% ===== Part 9: Visualize Weights =====
190 % You can now "visualize" what the neural network is learning by
191 % displaying the hidden units to see what features they are capturing in
192 % the data.
193
194 fprintf('\nVisualizing Neural Network... \n')
195
196 displayData(w1(:, 2:end));
197
198 fprintf('\nProgram paused. Press enter to continue.\n');
199 pause;
200
201 %% ===== Part 10: Implement Predict =====
202 % After training the neural network, we would like to use it to predict
203 % the labels. You will now implement the "predict" function to use the
204 % neural network to predict the labels of the training set. This lets
205 % you compute the training set accuracy.
206
207 pred = predict(w1, w2, X);
208
209 fprintf('\nTraining Set Accuracy: %f\n', mean(double(pred == y)) * 100);

```

```

1 function [C grad] = nnCostFunction(nn_params, ...
2                                     input_layer_size, ...
3                                     hidden_layer_size, ...
4                                     num_labels, ...
5                                     X, y, lambda)
6 %NNCOSTFUNCTION Implements the neural network cost function for a two layer
7 %neural network which performs classification
8
9 % Reshape nn_params back into the parameters weight1 and weight2, the weight
   matrices
10 % for our 2 layer neural network
11 w1 = reshape(nn_params(1:hidden_layer_size * (input_layer_size + 1)), ...
12              hidden_layer_size, (input_layer_size + 1));
13
14 w2 = reshape(nn_params((1 + (hidden_layer_size * (input_layer_size + 1))):end), ...
15              num_labels, (hidden_layer_size + 1));
16
17 % Setup some useful variables
18 m = size(X, 1);
19
20 % You need to return the following variables correctly
21 C = 0;
22 w1_grad = zeros(size(w1));
23 w2_grad = zeros(size(w2));
24
25 A1 = [ones(size(X,1),1) X];
26 Z2 = A1*w1';
27 A2 = sigmoid(Z2);
28 A2_0 = [ones(size(A2,1),1) A2];
29 Z3 = A2_0*w2';
30 A3 = sigmoid(Z3);
31
32 HX = A3;
33
34 % recode y to Y // TRICKY
35 I = eye(num_labels);
36 Y = zeros(m, num_labels);
37
38 for i=1:m
39     Y(i, :) = I(y(i), :);
40 end
41
42 penalty = (lambda / (2*m)) * (sum(sum(w1(:, 2:end).^2, 2)) + sum(sum(w2(:, 2:end).^2, 2)
   ));
43
44 C = (1/m) * sum(sum((-Y).*log(HX) - ((1-Y).*log(1-HX)))) + penalty;
45
46 % calculate sigmas //DELTAS
47 sigma3 = A3 - Y;
48 sigma2 = (sigma3*w2).*sigmoidGradient([ones(size(Z2, 1), 1) Z2]);
49 sigma2 = sigma2(:, 2:end);
50
51 % accumulate gradients
52 delta_1 = (sigma2'*A1);
53 delta_2 = (sigma3'*A2_0);
54

```

```

55 % calculate regularized gradient
56 p1 = (lambda/m)*[zeros(size(w1, 1), 1) w1(:, 2:end)];
57 p2 = (lambda/m)*[zeros(size(w2, 1), 1) w2(:, 2:end)];
58 w1_grad = delta_1./m + p1;
59 w2_grad = delta_2./m + p2;
60
61 % Unroll gradients
62 grad = [w1_grad(:) ; w2_grad(:)];
63
64 end

```

sigmoidGradient.m

```

1 function g = sigmoidGradient(z)
2 %SIGMOIDGRADIENT returns the gradient of the sigmoid function
3 %evaluated at z
4
5 g = zeros(size(z));
6 g = sigmoid(z).*(1-sigmoid(z));
7
8 end

```

randInitializeWeights.m

```

1 function W = randInitializeWeights(L_in, L_out)
2 %RANDINITIALIZEWEIGHTS Randomly initialize the weights of a layer with L_in
3 %incoming connections and L_out outgoing connections
4 % Note that W should be set to a matrix of size(L_out, 1 + L_in) as
5 % the first column of W handles the "bias" terms
6
7 W = zeros(L_out, 1 + L_in);
8 % Initialize W randomly so that we break the symmetry while
9 % training the neural network.
10
11 epsilon_init = 0.12;
12 W = rand(L_out, 1 + L_in) * 2 * epsilon_init - epsilon_init;
13
14 end

```

predict.m

```

1 function p = predict(w1, w2, X)
2 %PREDICT Predict the label of an input given a trained neural network
3 % p = PREDICT(w1, w2, X) outputs the predicted label of X given the
4 % trained weights of a neural network (w1, w2)
5
6 % Useful values
7 m = size(X, 1);
8 num_labels = size(w2, 1);
9
10 % You need to return the following variables correctly
11 p = zeros(size(X, 1), 1);
12
13 h1 = sigmoid([ones(m, 1) X] * w1');
14 h2 = sigmoid([ones(m, 1) h1] * w2');
15 [dummy, p] = max(h2, [], 2);
16
17 end

```