

```

1 %% Machine Learning
2 % Lab 8: Regularized Linear Regression and Bias-Variance
3 % — Water overflow —
4 %{
5 In the first half of the exercise, you will implement regularized linear
6 regression to predict the amount of water flowing out of a dam using the
7 change of water level in a reservoir. In the next half, you will go through
8 some diagnostics of debugging learning algorithms and examine the effects
9 of bias v.s. variance.
10 %}
11 %% Initialization
12 clear ; close all; clc
13
14 %% ===== Part 1: Loading and Visualizing Data =====
15 % We start the exercise by first loading and visualizing the dataset.
16 % The following code will load the dataset into your environment and plot
17 % the data.
18 %
19
20 % Load Training Data
21 fprintf('Loading and Visualizing Data ...\n')
22
23 % Load from ex5data1:
24 % You will have X, y, Xval, yval, Xtest, ytest in your environment
25 load ('ex5data1.mat');
26
27 % m = Number of examples
28 m = size(X, 1);
29
30 % Plot training data
31 figure(1)
32 plot(X, y, 'rx', 'MarkerSize', 10, 'LineWidth', 1.5);
33 xlabel('Change in water level (x)');
34 ylabel('Water flowing out of the dam (y)');
35
36 fprintf('Program paused. Press enter to continue.\n');
37 pause;
38
39 %% ===== Part 2: Regularized Linear Regression Cost =====
40 % You should now implement the cost function for regularized linear
41 % regression.
42 %
43
44 w = [1 ; 1];
45 C = linearRegCostFunction([ones(m, 1) X], y, w, 1);
46
47 fprintf(['Cost at theta = [1 ; 1]: %f '...
48         '\n(this value should be about 303.993192)\n'], C);
49
50 fprintf('Program paused. Press enter to continue.\n');
51 pause;
52
53 %% ===== Part 3: Regularized Linear Regression Gradient =====
54 % You should now implement the gradient for regularized linear
55 % regression.
56 %

```

```

57
58 w = [1 ; 1];
59 [C, grad] = linearRegCostFunction([ones(m, 1) X], y, w, 1);
60
61 fprintf(['Gradient at theta = [1 ; 1]:  [%f; %f] '...
62         '\n(this value should be about [-15.303016; 598.250744])\n'], ...
63         grad(1), grad(2));
64
65 fprintf('Program paused. Press enter to continue.\n');
66 pause;
67
68
69 %% ===== Part 4: Train Linear Regression =====
70 % Once you have implemented the cost and gradient correctly, the
71 % trainLinearReg function will use your cost function to train
72 % regularized linear regression.
73 %
74 % Write Up Note: The data is non-linear, so this will not give a great
75 %                 fit.
76 %
77
78 % Train linear regression with lambda = 0
79 lambda = 0;
80 [w] = trainLinearReg([ones(m, 1) X], y, lambda);
81
82 % Plot fit over the data
83 figure(1)
84 plot(X, y, 'rx', 'MarkerSize', 10, 'LineWidth', 1.5);
85 xlabel('Change in water level (x)');
86 ylabel('Water flowing out of the dam (y)');
87 hold on;
88 plot(X, [ones(m, 1) X]*w, '—', 'LineWidth', 2)
89 hold off;
90
91 fprintf('Program paused. Press enter to continue.\n');
92 pause;
93
94 %% ===== Part 5: Learning Curve for Linear Regression =====
95 % Next, you should implement the learningCurve function.
96 %
97 % Write Up Note: Since the model is underfitting the data, we expect to
98 %                 see a graph with "high bias"
99
100 lambda = 0;
101 [error_train, error_val] = ...
102     learningCurve([ones(m, 1) X], y, ...
103                 [ones(size(Xval, 1), 1) Xval], yval, ...
104                 lambda);
105
106 figure(2)
107 plot(1:m, error_train, 1:m, error_val);
108 title('Learning curve for linear regression')
109 legend('Train', 'Cross Validation')
110 xlabel('Number of training examples')
111 ylabel('Error')
112 axis([0 13 0 150])
113

```

```

114 fprintf('# Training Examples\tTrain Error\tCross Validation Error\n');
115 for i = 1:m
116     fprintf(' \t%d\t\t%f\t%f\n', i, error_train(i), error_val(i));
117 end
118
119 fprintf('Program paused. Press enter to continue.\n');
120 pause;
121
122 %% ===== Part 6: Feature Mapping for Polynomial Regression =====
123 % One solution to this is to use polynomial regression. You should now
124 % complete polyFeatures to map each example into its powers
125 %
126
127 p = 8;
128
129 % Map X onto Polynomial Features and Normalize
130 X_poly = polyFeatures(X, p);
131 [X_poly, mu, sigma] = featureNormalize(X_poly); % Normalize
132 X_poly = [ones(m, 1), X_poly]; % Add Ones
133
134 % Map X_poly_test and normalize (using mu and sigma)
135 X_poly_test = polyFeatures(Xtest, p);
136 X_poly_test = bsxfun(@minus, X_poly_test, mu);
137 X_poly_test = bsxfun(@rdivide, X_poly_test, sigma);
138 X_poly_test = [ones(size(X_poly_test, 1), 1), X_poly_test]; % Add Ones
139
140 % Map X_poly_val and normalize (using mu and sigma)
141 X_poly_val = polyFeatures(Xval, p);
142 X_poly_val = bsxfun(@minus, X_poly_val, mu);
143 X_poly_val = bsxfun(@rdivide, X_poly_val, sigma);
144 X_poly_val = [ones(size(X_poly_val, 1), 1), X_poly_val]; % Add Ones
145
146 fprintf('Normalized Training Example 1:\n');
147 fprintf(' %f \n', X_poly(1, :));
148
149 fprintf('\nProgram paused. Press enter to continue.\n');
150 pause;
151
152 %% ===== Part 7: Learning Curve for Polynomial Regression =====
153 % Now, you will get to experiment with polynomial regression with multiple
154 % values of lambda. The code below runs polynomial regression with
155 % lambda = 0. You should try running the code with different values of
156 % lambda to see how the fit and learning curve change.
157
158 lambda = 0;
159 [w] = trainLinearReg(X_poly, y, lambda);
160
161 % Plot training data and fit
162 figure(3);
163 plot(X, y, 'rx', 'MarkerSize', 10, 'LineWidth', 1.5);
164 plotFit(min(X), max(X), mu, sigma, w, p);
165 xlabel('Change in water level (x)');
166 ylabel('Water flowing out of the dam (y)');
167 title(sprintf('Polynomial Regression Fit (lambda = %f)', lambda));
168
169 figure(4);
170 [error_train, error_val] = ...

```

```

171     learningCurve(X_poly, y, X_poly_val, yval, lambda);
172 plot(1:m, error_train, 1:m, error_val);
173
174 title(sprintf('Polynomial Regression Learning Curve (lambda = %f)', lambda));
175 xlabel('Number of training examples')
176 ylabel('Error')
177 axis([0 13 0 100])
178 legend('Train', 'Cross Validation')
179
180 fprintf('Polynomial Regression (lambda = %f)\n\n', lambda);
181 fprintf('# Training Examples\tTrain Error\tCross Validation Error\n');
182 for i = 1:m
183     fprintf(' \t%d\t\t%f\t%f\n', i, error_train(i), error_val(i));
184 end
185
186 fprintf('Program paused. Press enter to continue.\n');
187 pause;
188
189 %% ===== Part 8: Validation for Selecting Lambda =====
190 % You will now implement validationCurve to test various values of
191 % lambda on a validation set. You will then use this to select the
192 % "best" lambda value.
193 %
194
195 [lambda_vec, error_train, error_val] = ...
196     validationCurve(X_poly, y, X_poly_val, yval);
197
198 %close all;
199 figure(5)
200 plot(lambda_vec, error_train, lambda_vec, error_val);
201 legend('Train', 'Cross Validation');
202 xlabel('lambda');
203 ylabel('Error');
204
205 fprintf('lambda\t\tTrain Error\tValidation Error\n');
206 for i = 1:length(lambda_vec)
207     fprintf(' %f\t%f\t%f\n', ...
208         lambda_vec(i), error_train(i), error_val(i));
209 end
210
211 fprintf('Program paused. Press enter to continue.\n');
212 pause;

```

linearRegCostFunction.m

```

1 function [C, grad] = linearRegCostFunction(X, y, w, lambda)
2 %LINEARREGCOSTFUNCTION Compute cost and gradient for regularized linear
3 %regression with multiple variables
4
5 % Initialize some useful values
6 m = length(y); % number of training examples
7
8 % You need to return the following variables correctly
9 C = 0;
10 grad = zeros(size(w));
11
12 w_0 = w;
13 w_0(1) = 0;
14 penalty = (lambda / (2*m))*sum(w_0.^2);
15 C = (1/(2*m))*(sum((X*w)-y).^2)+penalty;
16
17 grad = (1/m)*sum(((X*w)-y).*X)+(lambda/m)*w_0';
18
19 grad = grad(:);
20
21 end

```

trainLinearReg.m

```

1 function [w] = trainLinearReg(X, y, lambda)
2 %TRAINLINEARREG Trains linear regression given a dataset (X, y) and a
3 %regularization parameter lambda
4
5 % Initialize weights
6 init_w = zeros(size(X, 2), 1);
7
8 % Create "short hand" for the cost function to be minimized
9 costFunction = @(t) linearRegCostFunction(X, y, t, lambda);
10
11 % Now, costFunction is a function that takes in only one argument
12 options = optimset('MaxIter', 200, 'GradObj', 'on');
13
14 % Minimize using fmincg
15 w = fmincg(costFunction, init_w, options);
16
17 end

```

learningCurve.m

```

1 function [error_train, error_val] = ...
2     learningCurve(X, y, Xval, yval, lambda)
3 %LEARNINGCURVE Generates the train and cross validation set errors needed
4 %to plot a learning curve
5
6 % Number of training examples
7 m = size(X, 1);
8
9 % You need to return these values correctly
10 error_train = zeros(m, 1);
11 error_val = zeros(m, 1);

```

```

12
13 for i= 1:m
14     w          = trainLinearReg(X(1:i,:), y(1:i), lambda);
15     error_train(i) = linearRegCostFunction(X(1:i,:), y(1:i) , w, 0);
16     error_val(i)   = linearRegCostFunction(Xval, yval, w, 0);
17 end
18
19 end

```

polyFeatures.m

```

1 function [X_poly] = polyFeatures(X, p)
2 %POLYFEATURES Maps X (1D vector) into the p-th power
3
4 % You need to return the following variables correctly.
5 X_poly = zeros(numel(X), p);
6
7 for i=1:p
8     X_poly(:,i)=X.^i;
9 end
10
11 end

```

featureNormalize.m

```

1 function [X_norm, mu, sigma] = featureNormalize(X)
2 %FEATURENORMALIZE Normalizes the features in X
3 %{
4 bsxfun  Binary Singleton Expansion Function
5     C = bsxfun(FUNC,A,B) applies the element-by-element binary operation
6     specified by the function handle FUNC to arrays A and B, with implicit
7     expansion enabled.
8 %}
9 mu = mean(X);
10 X_norm = bsxfun(@minus, X, mu);
11
12 sigma = std(X_norm);
13 X_norm = bsxfun(@rdivide, X_norm, sigma);
14
15 end

```

validationCurve.m

```
1 function [lambda_vec, error_train, error_val] = ...
2     validationCurve(X, y, Xval, yval)
3 %VALIDATIONCURVE Generate the train and validation errors needed to
4 %plot a validation curve that we can use to select lambda
5
6 % Selected values of lambda (you should not change this)
7 lambda_vec = [0 0.001 0.003 0.01 0.03 0.1 0.3 1 3 10]';
8
9 % You need to return these variables correctly.
10 error_train = zeros(length(lambda_vec), 1);
11 error_val = zeros(length(lambda_vec), 1);
12
13 for i = 1:length(lambda_vec)
14     lambda = lambda_vec(i);
15     theta = trainLinearReg(X, y, lambda);
16     error_train(i) = linearRegCostFunction(X, y, theta, 0);
17     error_val(i) = linearRegCostFunction(Xval, yval, theta, 0);
18 end
19
20 end
```